



# Image Captioning

Group 6 – Yunhong Yang, Ziyue Li, Yutao Chen, Xiaodan Lu

# Table of Contents

**01 - Introduction**

**02 - Motivations**

**03 - Datasets**

**04 - Methodology**

**05 - Results**

**06 - Conclusion**

# Introduction



**What do you see in this picture?**

----- “A young girl is eating broccoli in cream sauce.”

# Introduction

---

- Image Captioning is the task of describing the content of an image in words.
- Image captioning is an intricate endeavor bridging computer vision and natural language processing (NLP).
- In this, the input to the model is an image and the model's output is a caption generated in natural language processing.

# Motivations

The first thing to ponder is the significance of the image captioning problem statement. Image captioning has a huge amount of application. Let's explore a few noteworthy examples:

- Aid to the Blind
- Autonomous vehicles
- Google Image Search

# MS COCO Dataset

The MS COCO (Microsoft Common Objects in Context) dataset is a comprehensive collection encompassing object detection, segmentation, key-point detection, and captioning tasks. This extensive dataset comprises more than 200,000 images.



“A man wearing earphones doing a trick on a skateboard ramp.”



“There is a surfer wearing a body suit riding a wave.”



“The teenagers are standing together on the sidewalk.”



# Data Preprocessing

## Data Cleaning

- Converted text to lowercase.
- Remove extra white spaces.
- Removed punctuation and special characters.
- Added Start and End Tokens.

## Data Splitting

- Train Set: 80%
- Validation Set: 10%
- Test Set: 10%
- Sizes: Training (7228), Validation (903), Test (904).

## Tokenization & Padding

- Map tokens to integers.
- Utilized padding to ensure consistent sequence length.

## Resizing Images

- Resized each image to the specified dimensions to ensure consistent input size for neural networks.

# Methodology

- **Encoder**
- **Embedding Layer**
- **Decoder**
- **Evaluation Metrics**



# Encoder

## InceptionV3 Model

- Inception v3 is a convolutional neural network for assisting in image analysis and object detection, and got its start as a module for GoogLeNet.
- The InceptionV3 model is loaded with pre-trained weights, and its layers are frozen to prevent further training.

# RepeatVector

- Why: Single image encoding may not capture enough dynamic details
- Purpose: Replicate image encoding
- Effect: Ensures consistent image context for every word
- Benefits:
  - Enhanced model understanding.
  - Improved generation of contextually relevant captions.

Source: <https://towardsdatascience.com/step-by-step-understanding-lstm-autoencoder-layers-ffab055b6352>



# Embedding

## How it works

- Mapping: Converts words to vectors in a continuous space.
- Contextualization: Embeddings capture semantic relationships.
- Dropout: Introduces regularization, preventing overfitting.

Source: <https://www.codingninjas.com/studio/library/embedding-layers-in-keras>

# Encoding & Embedding Process

Encoding → Encoding → Encoding → Embedding → Concatenating

Preprocess input and add dense layer for image encoding with ReLU activation

Flatten the encoding to be used as input in subsequent layers

Repeat the flattened encoding for each time step in the caption

Embedding layer for the teacher forcing input and apply dropout

Concatenate the flattened encoding and embedded teacher forcing input

# LSTM: Decoder

Why?

- Decoding the fixed-length vector and outputting the predicted sequence

In the encoder part: only one vector in the last time step and neglecting all the others

In the decoder part: an output vector at every time step so the Dense layer can make a prediction.



# Dense Layer

Why?

- The decoder is just a language model conditioned on the initial states.

last step:  
Predict the Caption.

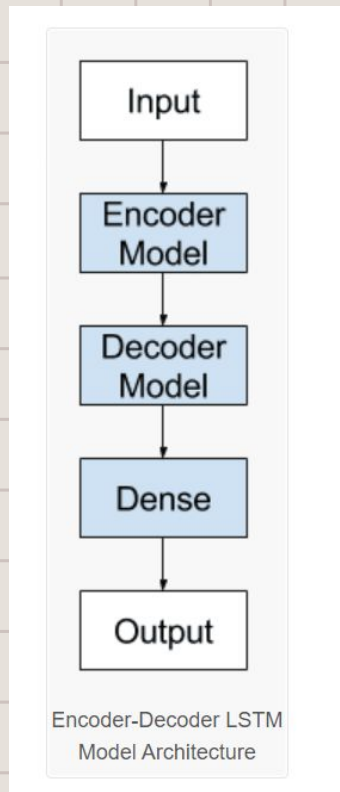
The number of units is the shape of the output vector



Source:

<https://towardsdatascience.com/how-to-build-an-encoder-decoder-translation-model-using-lstm-with-python-and-keras-a31e9d864b9b>

# Summary



Input: Input Text data(Caption) and Image Data

Encoder: InceptionV3 Model as encoder to extract features from the image. processes an input sequence and generates an encoded state

Decoder: LSTM as decoder. uses the encoded state to produce an output sequence.

Dense: Predict the Caption

Output: A sentence of the image caption

# Loss Function

```
# Loss function
def sparse_it_up(y_true, y_preds):
    # Cast true labels to integers
    y_true = tf.cast(y_true, tf.int32)
    # Create a mask for non-zero elements in true labels
    mask = tf.math.logical_not(tf.math.equal(y_true, 0))
    mask = tf.cast(mask, dtype=tf.float32)
    # Convert true labels to one-hot encoding
    y_true_one_hot = tf.one_hot(y_true, words + 1)
    # Calculate categorical cross-entropy with masking
    loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y_true_one_hot, y_preds) * mask)
    return loss
```

Cross entropy serves as the loss function for model training.

1. converts the `y_true` tensor to integer type, which is suitable for handling class indices.
2. ignore loss contributions from padded values (where `y_true` equals 0)
3. The `y_true` tensor is one-hot encoded. The `words + 1` represents the number of classes, and each class is assigned a unique one-hot vector.
4. used to compute the cross-entropy loss between the predicted logits (`y_preds`) and the true labels (`y_true_one_hot`).



# Evaluation Metrics

Metrics used: BLEU, METEOR and ROUGE scores

Metric	Problem Solved	Method
ROUGE	Text summarization	Measures overlap of N-grams and LCS between system-generated and reference summaries
BLEU	Machine translation	Measures N-gram precision between candidate and reference translations
METEOR	Machine translation	Harmonic mean of unigram precision and recall, with a penalty for length mismatches



[Image source](#)

# Results



# Result

## Loss Function




# Result

Average BLEU: 0.0941

Average METEOR: 0.3059

Average ROUGE: 0.3090



# Conclusion





# Conclusion

**Label:** a man riding a skateboard down a sidewalk

**Predicted:** a man riding a skateboard while a skateboard



**Label:** a man on a snowboard who is performing a jump

**Predicted:** a person riding a snowboard with is skiing a snowboard



**Label:** a child laying on a bed in a room

**Predicted:** a small is on a bed with a bed



# Limitation & Future Thoughts

## Limitation

Only 10,000 datasets about “person” were selected in coco dataset.

## Future Thoughts

Exploring more diverse datasets  
Improving its robustness and versatility



Image source





# Thank you!

Q & A

# References:

- BM, Nechu. "How to Build an Encoder Decoder Translation Model Using LSTM with Python and Keras." *Medium*, Towards Data Science, 21 Oct. 2020, [towardsdatascience.com/how-to-build-an-encoder-decoder-translation-model-using-lstm-with-python-and-keras-a31e9d864b9b](https://towardsdatascience.com/how-to-build-an-encoder-decoder-translation-model-using-lstm-with-python-and-keras-a31e9d864b9b).
- Dhote, Pradeep. "Seq2Seq-Encoder-Decoder-LSTM-Model." *Medium*, Medium, 20 Aug. 2020, [pradeep-dhote9.medium.com/seq2seq-encoder-decoder-lstm-model-1a1c9a43bbac](https://pradeep-dhote9.medium.com/seq2seq-encoder-decoder-lstm-model-1a1c9a43bbac).
- "Rouge Score." *TorchMetrics*, [torchmetrics.readthedocs.io/en/stable/text/rouge\\_score.html](https://torchmetrics.readthedocs.io/en/stable/text/rouge_score.html). Accessed 30 Nov. 2023.
- Awsaf. "Coco 2017 Dataset." *Kaggle*, 4 Sept. 2020, [www.kaggle.com/datasets/awsaf49/coco-2017-dataset](https://www.kaggle.com/datasets/awsaf49/coco-2017-dataset).
- Mansuy, Raphael. "Evaluating the Performance of Natural Language Processing (NLP) Models Can Be Challenging." *Medium*, 27 Nov. 2023, [ai.plainenglish.io/evaluating-the-performance-of-natural-language-processing-nlp-models-can-be-challenging-ce6f62c07c35](https://ai.plainenglish.io/evaluating-the-performance-of-natural-language-processing-nlp-models-can-be-challenging-ce6f62c07c35).